Cognify Project Test Document

Introduction to Cognify Testing

This is a sample document created specifically for testing the Cognify RAG (Retrieval-Augmented Generation) application. Cognify is designed to ingest documents, chunk them, embed them using OpenAI's embedding models, and store them in Qdrant for semantic search.

Features of Cognify:

1. Document Upload: Users can upload PDF documents up to 50MB.

2. Vector Storage: Uses Qdrant for high-performance vector search.

3. LLM Integration: Uses OpenAI's GPT models to generate answers based on retrieved context.

4. Streaming Responses: Provides real-time typing effect via Server-Sent Events (SSE).

System Architecture:

The system uses FastAPI for the backend API, SQLAlchemy with PostgreSQL for relational metadata, Redis for caching and rate limiting, and Qdrant for the vector database. When a user uploads a document, the document is parsed, chunked according to the configured chunk size (1000) and overlap (200), and then embedded.

Testing Scenarios:

- Upload this document via the API.

- Query: What is the maximum file size for upload?

- Query: Which vector database does Cognify use?

- Query: What technologies are used in the system architecture?

End of Test Document.